

Two Schedulers in PA1

Implementing Schedulers on XINU

- Solving the starvation problem
- Suggested background knowledge
 - The default scheduler (resched.c) and scheduling points (where it gets invoked)
 - Process management
 - Creation, suspension, resumption, sleep, wakeup, termination, priority, etc. ...
 - Process queues (ready queue, clock queue) and their management (insert, dequeue, etc. ...)

Exponential Distribution Scheduler

- Ensuring fairness by probability distribution and preemptions
- Important implementation points
 - The random number generator $X \sim \text{Exp}(\lambda = 0.1)$
 - Scheduling a **runnable** process (i.e., processes from the ready queue) whose priority is the lowest one above the outcome of X (or the highest if none is above)
 - Round robin on processes with the same priority
- Examining scheduling results against
 - Probability density function: $f(x; \lambda) = \lambda e^{-\lambda x}$
 - Cumulative distribution function: $F(x; \lambda) = 1 - e^{-\lambda x}$

Linux-like Scheduler

- Ensuring fairness by predetermined time quota and epochs
- Important concepts
 - **Epoch**: a turn that every runnable process is guaranteed a non-preemptive chance to run up to its predetermined time quantum
 - **Time quantum**: CPU ticks specifying how long a process can run within an epoch
 - **Base priority**: a static priority specified by `create()` or `chprio()`
 - **Goodness value**: a dynamic priority indicating when a process can be scheduled to run within an epoch

Determining Time Quantum and Goodness Value

- Time quantum

- Base priority + $0.5 * \text{unused quantum from the previous epoch}$
- Determined at the beginning of each epoch
- Computed for all processes including ones not in the ready queue

- Goodness value

- Base priority + unused time quantum in the current epoch
- 0 if time quantum is used up
- Updated along with the unused time quantum whenever the scheduler is invoked

Scheduling Policies

- Starting a new epoch when all runnable processes have used up their time
- Scheduling the process with the highest goodness
 - Using the round-robin strategy if multiple processes have the same goodness
 - Be aware of the NULL process
- New processes and priority changes not in effect until the next epoch

Expected Results

- In each epoch, every runnable process will eventually be scheduled to run without being preempted until:
 - It uses up its time quantum
 - It yields (e.g., via invoking sleep) with unused time quantum
- A process with unused time quantum will be rescheduled if it becomes runnable again within the same epoch
- In other cases, half of the unused time quantum is carried to the next epoch

Other Implementation Notes

- Implementing the following functions to allow scheduler switching
 - void setschedclass (int sched_class)
 - Setting a scheduler type
 - sched_class: EXPDISTSCHEDED = 1, LINUXSCHEDED = 2
 - int getschedclass(): getting the current scheduler type
- Making your program compatible with the testmain.c that we provide